

Table 4 describes the Gst-nvinfer plugin's Gst properties.

Table 4. Gst-nvinfer plugin, Gst properties

Property	Meaning	Type and Range	Example Notes
config-file-path	Absolute pathname of configuration file for the Gst-nvinfer element	String	config-file-path= config_infer_primary.txt
process-mode	Infer Processing Mode 1=Primary Mode 2=Secondary Mode	Integer, 1 or 2	process-mode=1
unique-id	Unique ID identifying metadata generated by this GIE	Integer, 0 to 4,294,967,295	unique-id=1
infer-on-gie-id	See operate-on-gie-id in the configuration file table	Integer, 0 to 4,294,967,295	infer-on-gie-id=1
infer-on-class-ids	See operate-on-class-ids in the configuration file table	An array of colon-separated integers (class-ids)	infer-on-class-ids=1:2:4
model-engine-file	Absolute pathname of the pre-generated serialized engine file for the mode	String	model-engine-file= model_b1_fp32.engine
batch-size	Number of frames/objects to be inferred together in a batch	Integer, 1 - 4,294,967,295	batch-size=4
Interval	Number of consecutive batches to be skipped for inference	Integer, 0 to 32	interval=0
gpu-id	Device ID of GPU to use for pre-processing/inference (dGPU only)	Integer, 0-4,294,967,295	gpu-id=1
raw-output-file-write	Pathname of raw inference output file	Boolean	raw-output-file-write=1
raw-output-generated-callback	Pointer to the raw output generated callback function	Pointer	<i>Cannot be set through gst-launch</i>
raw-output-generated-userdata	Pointer to user data to be supplied with raw-output-generated-callback	Pointer	<i>Cannot be set through gst-launch</i>
output-tensor-meta	Indicates whether to attach tensor outputs as meta on GstBuffer.	Boolean	output-tensor-meta=0

2.1.5 Tensor Metadata

The Gst-nvinfer plugin can attach raw output tensor data generated by a TensorRT inference engine as metadata. It is added as an `NvDsInferTensorMeta` in the

`frame_user_meta_list` member of `NvDsFrameMeta` for primary (full-frame) mode, or in the `obj_user_meta_list` member of `NvDsObjectMeta` for secondary (object) mode.

To read or parse inference raw tensor data of output layers

1. Enable property `output-tensor-meta`, or enable the same-named attribute in the configuration file for the `Gst-nvinfer` plugin.
2. When operating as primary GIE, `NvDsInferTensorMeta` is attached to each frame's (each `NvDsFrameMeta` object's) `frame_user_meta_list`. When operating as secondary GIE, `NvDsInferTensorMeta` is attached to each `NvDsObjectMeta` object's `obj_user_meta_list`.

Metadata attached by `Gst-nvinfer` can be accessed in a `GStreamer` pad probe attached downstream from the `Gst-nvinfer` instance.

3. The `NvDsInferTensorMeta` object's metadata type is set to `NVDSINFER_TENSOR_OUTPUT_META`. To get this metadata you must iterate over the `NvDsUserMeta` user metadata objects in the list referenced by `frame_user_meta_list` or `obj_user_meta_list`.

For more information about `Gst-infer` tensor metadata usage, see the source code in `sources/apps/sample_apps/deepstream_infer_tensor_meta-test.cpp`, provided in the DeepStream SDK samples.

2.1.6 Segmentation Metadata

The `Gst-nvinfer` plugin attaches the output of the segmentation model as user meta in an instance of `NvDsInferSegmentationMeta` with `meta_type` set to `NVDSINFER_SEGMENTATION_META`. The user meta is added to the `frame_user_meta_list` member of `NvDsFrameMeta` for primary (full-frame) mode, or the `obj_user_meta_list` member of `NvDsObjectMeta` for secondary (object) mode.

For guidance on how to access user metadata, see [User/Custom Metadata Addition inside NvDsBatchMeta](#), and [Tensor Metadata](#), above.

2.2 GST-NVTRACKER

This plugin tracks detected objects and gives each new object a unique ID.

The plugin adapts a low-level tracker library to the pipeline. It supports any low-level library that implements the low-level API, including the three reference implementations, the `NvDCF`, `KLT`, and `IOU` trackers.

As part of this API, the plugin queries the low-level library for capabilities and requirements concerning input format and memory type. It then converts input buffers into the format requested by the low-level library. For example, the KLT tracker uses Luma-only format; NvDCF uses NV12 or RGBA; and IOU requires no buffer at all.

The low-level capabilities also include support for batch processing across multiple input streams. Batch processing is typically more efficient than processing each stream independently. If a low-level library supports batch processing, that is the preferred mode of operation. However, this preference can be overridden with the `enable-batch-process` configuration option if the low-level library supports both batch and per-stream modes.

The plugin accepts NV12/RGBA data from the upstream component and scales (converts) the input buffer to a buffer in the format required by the low-level library, with tracker width and height. (Tracker width and height must be specified in the configuration file's `[tracker]` section.)

The low-level tracker library is selected via the `ll-lib-file` configuration option in the tracker configuration section. The selected low-level library may also require its own configuration file, which can be specified via the `ll-config-file` option.

The three reference low level libraries support different algorithms:

- ▶ The KLT tracker uses a CPU-based implementation of the Kanade Lucas Tomasi (KLT) tracker algorithm. This library requires no configuration file.
- ▶ The Intersection of Union (IOU) tracker uses the intersection of the detector's bounding boxes across frames to determine the object's unique ID. This library takes an optional configuration file.
- ▶ The Nv-adapted Discriminative Correlation Filter (NvDCF) tracker uses a correlation filter-based online discriminative learning algorithm, coupled with a Hungarian

algorithm for data association in multi-object tracking. This library accepts an optional configuration file

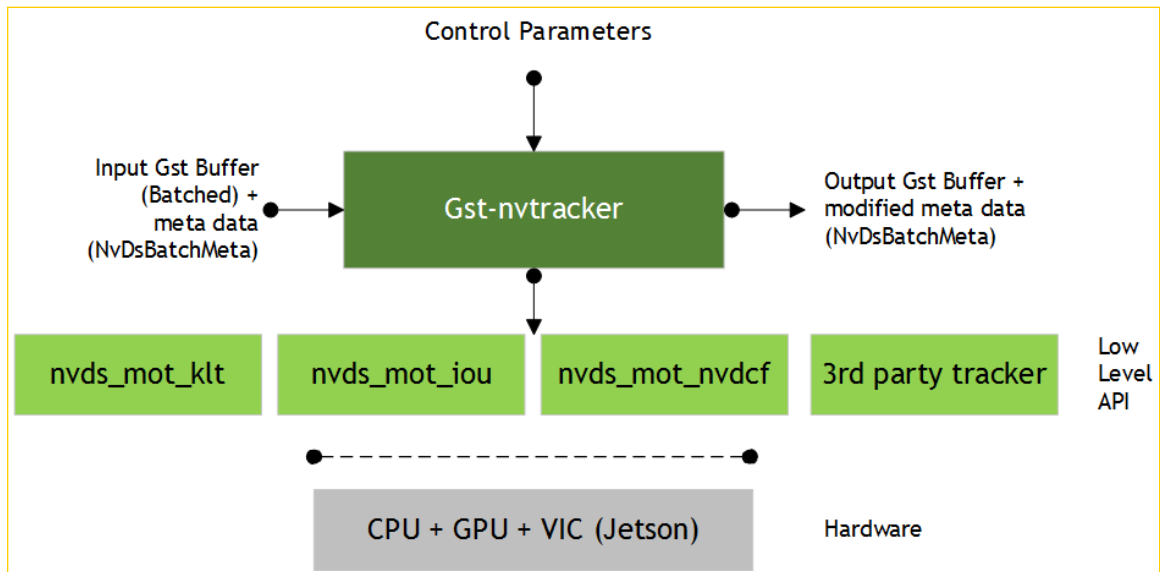


Figure 2. Gst-nvtracker inputs and outputs

2.2.1 Inputs and Outputs

This section summarizes the inputs, outputs, and communication facilities of the Gst-nvtracker plugin.

► Inputs

- Gst Buffer (batched)
- NvDsBatchMeta

Formats supported are NV12 and RGBA.

► Control parameters

- tracker-width
- tracker-height
- gpu-id (for dGPU only)
- ll-lib-file
- ll-config-file
- enable-batch-process

► Output

- Gst Buffer (provided as an input)
- NvDsBatchMeta (Updated by Gst-nvtracker with tracked object coordinates and object IDs)

2.2.2 Features

Table 5 summarizes the features of the plugin.

Table 5. Features of the Gst-nvtracker plugin

Feature	Description	Release
Configurable tracker width/height	Frames are internally scaled to specified resolution for tracking	DS2.0
Multi-stream CPU/GPU tracker	Supports tracking on batched buffers consisting of frames from different sources	DS2.0
NV12 Input	—	DS2.0
RGBA Input	—	DS 3.0
Allows low FPS tracking	IOU tracker	DS 3.0
Configurable GPU device	User can select GPU for internal scaling/color format conversions and tracking	DS2.0
Dynamic addition/deletion of sources at runtime	Supports tracking on new sources added at runtime and cleanup of resources when sources are removed	DS 3.0
Support for user's choice of low-level library	Dynamically loads user selected low-level library	DS 4.0
Support for batch processing	Supports sending frames from multiple input streams to the low-level library as a batch if the low-level library advertises capability to handle that	DS 4.0
Support for multiple buffer formats as input to low-level library	Converts input buffer to formats requested by the low-level library, for up to 4 formats per frame	DS 4.0

2.2.3 Gst Properties

Table 6 describes the Gst properties of the Gst-nvtracker plugin.

Table 6. Gst-nvtracker plugin, Gst Properties

Property	Meaning	Type and Range	Example Notes
tracker-width	Frame width at which the tracker is to operate, in pixels.	Integer, 0 to 4,294,967,295	tracker-width=640
tracker-height	Frame height at which the tracker is to operate, in pixels.	Integer, 0 to 4,294,967,295	tracker-height=368
ll-lib-file	Pathname of the low-level tracker library to be loaded by Gst-nvtracker.	String	ll-lib-file=/opt/nvidia/deepstream/libnvs_nvdcf.so

Property	Meaning	Type and Range	Example Notes
ll-config-file	Configuration file for the low-level library if needed.	Path to configuration file	ll-config-file=/opt/nvidia/deepstream/tracker_config.yml
gpu-id	ID of the GPU on which device/unified memory is to be allocated, and with which buffer copy/scaling is to be done. (dGPU only.)	Integer, 0 to 4,294,967,295	gpu-id=1
enable-batch-process	Enables/disables batch processing mode. Only effective if the low-level library supports both batch and per-stream processing. (Optional.)	Boolean	enable-batch-process=1

2.2.4 Custom Low-Level Library

To write a custom low-level tracker library, implement the API defined in `sources/includes/nvdstracker.h`. Parts of the API refer to `sources/includes/nvbufsurface.h`.

The names of API functions and data structures are prefixed with `NvMOT`, which stands for *NVIDIA Multi-Object Tracker*.

This is the general flow of the API from a low-level library perspective:

1. The first required function is:

```
NvMOTStatus NvMOT_Query(
    uint16_t customConfigFilePathSize,
    char* pCustomConfigFilePath,
    NvMOTQuery *pQuery
);
```

The plugin uses this function to query the low-level library's capabilities and requirements before it starts any processing sessions (contexts) with the library. Queried properties include the input frame memory format (e.g., RGBA or NV12), memory type (e.g., NVIDIA® CUDA® device or CPU mapped NVMM), and support for batch processing.

The plugin performs this query once, and its results apply to all contexts established with the low-level library. If a low-level library configuration file is specified, it is provided in the query for the library to consult.

The query reply structure `NvMOTQuery` contains the following fields:

- `NvMOTCompute computeConfig`: Reports compute targets supported by the library. The plugin currently only echoes the reported value when initiating a context.
- `uint8_t numTransforms`: The number of color formats required by the low-level library. The valid range for this field is 0 to `NVMOT_MAX_TRANSFORMS`. Set this to 0 if the library does not require any visual data. Note that 0 does not mean that untransformed data will be passed to the library.
- `NvBufSurfaceColorFormat colorFormats[NVMOT_MAX_TRANSFORMS]`: The list of color formats required by the low-level library. Only the first `numTransforms` entries are valid.
- `NvBufSurfaceMemType memType`: Memory type for the transform buffers. The plugin allocates buffers of this type to store color and scale-converted frames, and the buffers are passed to the low-level library for each frame. Note that support is currently limited to the following types:

dGPU: `NVBUF_MEM_CUDA_PINNED`

`NVBUF_MEM_CUDA_UNIFIED`

Jetson: `NVBUF_MEM_SURFACE_ARRAY`

- `bool supportBatchProcessing`: True if the low-library support batch processing across multiple streams; otherwise false.

2. After the query, and before any frames arrive, the plugin must initialize a context with the low-level library by calling:

```
NvMOTStatus NvMOT_Init(
    NvMOTConfig *pConfigIn,
    NvMOTContextHandle *pContextHandle,
    NvMOTConfigResponse *pConfigResponse
);
```

The context handle is opaque outside the low-level library. In batch processing mode, the plugin requests a single context for all input streams. In per-stream processing mode, the plugin makes this call for each input stream so that each stream has its own context.

This call includes a configuration request for the context. The low-level library has an opportunity to:

- Review the configuration, and create a context only if the request is accepted. If any part of the configuration request is rejected, no context is created, and the return status must be set to `NvMOTStatus_Error`. The `pConfigResponse` field can optionally contain status for specific configuration items.
- Pre-allocate resources based on the configuration.

Note:

- In the `NvMOTMiscConfig` structure, the `logMsg` field is currently unsupported and uninitialized.
- The `customConfigFilePath` pointer is only valid during the call.

3. Once a context is initialized, the plugin sends frame data along with detected object bounding boxes to the low-level library each time it receives such data from upstream. It always presents the data as a batch of frames, although the batch contains only a single frame in per-stream processing contexts. Each batch is guaranteed to contain at most one frame from each stream.

The function call for this processing is:

```
NvMOTStatus NvMOT_Process(NvMOTContextHandle contextHandle,  
    NvMOTProcessParams *pParams,  
    NvMOTTrackedObjBatch *pTrackedObjectsBatch  
);
```

Where:

- `pParams` is a pointer to the input batch of frames to process. The structure contains a list of one or more frames, with at most one frame from each stream. No two frame entries have the same `streamID`. Each entry of frame data contains a list of one or more buffers in the color formats required by the low-level library, as well as a list of object descriptors for the frame. Most libraries require at most one color format.
- `pTrackedObjectsBatch` is a pointer to the output batch of object descriptors. It is pre-populated with a value for `numFilled`, the number of frames included in the input parameters.

If a frame has no output object descriptors, it is still counted in `numFilled` and is represented with an empty list entry (`NvMOTTrackedObjList`). An empty list entry has the correct `streamID` set and `numFilled` set to 0.

Note:

The output object descriptor `NvMOTTrackedObj` contains a pointer to the associated input object, `associatedObjectIn`. You must set this to the associated input object only for the frame where the input object is passed in. For example:

- Frame 0: `NvMOTObjToTrack X` is passed in. The tracker assigns it ID 1, and the output object `associatedObjectIn` points to X.
- Frame 1: Inference is skipped, so there is no input object. The tracker finds object 1, and the output object `associatedObjectIn` points to NULL.
- Frame 2: `NvMOTObjToTrack Y` is passed in. The tracker identifies it as object 1. The output object 1 has `associatedObjectIn` pointing to Y.

4. When all processing is complete, the plugin calls this function to clean up the context:

```
void NvMOT_DeInit(NvMOTContextHandle contextHandle);
```

2.2.5 Low-Level Tracker Library Comparisons and Tradeoffs

DeepStream 4.0 provides three low-level tracker libraries which have different resource requirements and performance characteristics, in terms of accuracy, robustness, and efficiency, allowing you to choose the best tracker based on your use case. See the following table for comparison.

Table 7. Tracker library comparison

Tracker	Computational Load		Pros	Cons	Best Use Cases
	GPU	CPU			
IOU	X	Very Low	Light-weight	No visual features for matching, so prone to frequent tracker ID switches and failures. Not suitable for fast moving scene.	Objects are sparsely located, with distinct sizes. Detector is expected to run every frame or very frequently (ex. every alternate frame).
KLT	X	High	Works reasonably well for simple scenes	High CPU utilization. Susceptible to change in the visual appearance due to noise and perturbations, such as shadow, non-rigid deformation, out-of-plane rotation, and partial occlusion. Cannot work on objects with low textures.	Objects with strong textures and simpler background. Ideal for high CPU resource availability.
NvDCF	Medium	Low	Highly robust against partial occlusions, shadow, and other transient visual changes. Less frequent ID switches.	Slower than KLT and IOU due to increased computational complexity. Reduces the total number of streams processed.	Multi-object, complex scenes with partial occlusion.

2.2.6 NvDCF Low-Level Tracker

NvDCF is a reference implementation of the custom low-level tracker library that supports multi-stream, multi-object tracking in a batch mode using a discriminative correlation filter (DCF) based approach for visual object tracking and a Hungarian algorithm for data association.

NvDCF preallocates memory during initialization based on:

- ▶ The number of streams to be processed
- ▶ The maximum number of objects to be tracked per stream (denoted as `maxTargetsPerStream` in a configuration file for the NvDCF low-level library, `tracker_config.yml`)

Once the number of objects being tracked reaches the configured maximum value, any new objects will be discarded until resources for some existing tracked objects are released. Note that the number of objects being tracked includes objects that are tracked in Shadow Mode (described below). Therefore, NVIDIA recommends that you make `maxTargetsPerStream` large enough to accommodate the maximum number of objects of interest that may appear in a frame, as well as the past objects that may be tracked in shadow mode. Also, note that GPU memory usage by NvDCF is linearly proportional to the total number of objects being tracked, which is $(\text{number of video streams}) \times (\text{maxTargetsPerStream})$.

DCF-based trackers typically apply an exponential moving average for temporal consistency when the optimal correlation filter is created and updated. The learning rate for this moving average can be configured as `filterLr`. The standard deviation for Gaussian for desired response when creating an optimal DCF filter can also be configured as `gaussianSigma`.

DCF-based trackers also define a search region around the detected target location large enough for the same target to be detected in the search region in the next frame. The `SearchRegionPaddingScale` property determines the size of the search region as a multiple of the target's bounding box size. For example, with `SearchRegionPaddingScale: 3`, the size of the search region would be:

$$\text{searchregionwidth} = w + 3 * (w * h)^{1/2}$$

$$\text{searchregionheight} = h + 3 * (w * h)^{1/2}$$

Where `w` and `h` are the width and height of the target's bounding box.

Once the search region is defined for each target, the image patches for the search regions are cropped and scaled to a predefined feature image size, then the visual features are extracted. The `featureImgSizeLevel` property defines the size of the feature image. A lower value of `featureImgSizeLevel` causes NvDCF to use a smaller feature size, increasing GPU performance at the cost of accuracy and robustness.

Consider the relationship between `featureImgSizeLevel` and `SearchRegionPaddingScale` when configuring the parameters. If `SearchRegionPaddingScale` is increased while `featureImgSizeLevel` is fixed, the number of pixels corresponding to the target in the feature images is effectively decreased.

The `minDetectorConfidence` property sets confidence level below which object detection results are filtered out.

To achieve robust tracking, NvDCF employs two approaches to handling false alarms from PGIE detectors: **late activation** for handling false positives and **shadow tracking** for false negatives. Whenever a new object is detected a new tracker is instantiated in temporary mode. It must be activated to be considered as a valid target. Before it is activated it undergoes a probationary period, defined by `probationAge`, in temporary mode. If the object is not detected in more than `earlyTerminationAge` consecutive frames during the period, the tracker is terminated early.

Once the tracker for an object is activated, it is put into inactive mode only when (1) no matching detector input is found during the data association, or (2) the tracker confidence falls below a threshold defined by `minTrackerConfidence`. The per-object tracker will be put into active mode again if a matching detector input is found. The length of period during which a per-object tracker is in inactive mode is called the **shadow tracking age**; if it reaches the threshold defined by `maxShadowTrackingAge`, the tracker is terminated. If the bounding box of an object being tracked goes partially out of the image frame and so its visibility falls below a predefined threshold defined by `minVisibiilty4Tracking`, the tracker is also terminated.

Note that `probationAge` is counted against a clock that is incremented at every frame, while `maxShadowTrackingAge` and `earlyTerminationAge` are counted against a clock incremented only when the shadow tracking age is incremented. When the PGIE detector runs on every frame (i.e., `interval=0` in the `[primary-gie]` section of the `deepstream-app` configuration file), for example, all the ages are incremented based on the same clock. If the PGIE detector runs on every other frame (i.e., `interval` is set to 1 in `[primary-gie]`) and `probationAge: 12`, it will yield almost the same results as `interval=0` with `probationAge: 24`, because `shadowTrackingAge` would be incremented at a half speed compared to the case with PGIE `interval=0`.

Table 8 summaries the configuration parameters for an NvDCF low-level tracker.

Table 8. NvDCF low-level tracker, configuration properties

Property	Meaning	Type and Range	Example Notes
<code>maxTargetsPerStream</code>	Max number of targets to track per stream	Integer, 0 to 65535	<code>maxTargetsPerStream: 30</code>

Property	Meaning	Type and Range	Example Notes
filterLr	Learning rate for DCF filter in exponential moving average	Float, 0.0 to 1.0	filterLr: 0.11
gaussianSigma	Standard deviation for Gaussian for desired response	Float, >0.0	gaussianSigma: 0.75
minDetectorConfidence	Minimum detector confidence for a valid object	Float, -inf to inf	minDetectorConfidence: 0.0
minTrackerConfidence	Minimum detector confidence for a valid target	Float, 0.0 to 1.0	minTrackerConfidence: 0.6
featureImgSizeLevel	Size of a feature image	Integer, 1 to 5	featureImgSizeLevel : 1
SearchRegionPaddingScale	Search region size	Integer, 1 to 3	SearchRegionPaddingScale: 3
maxShadowTrackingAge	Maximum length of shadow tracking	Integer, ≥0	maxShadowTrackingAge: 9
probationAge	Length of probationary period	Integer, ≥0	probationAge: 12
earlyTerminationAge	Early termination age	Integer, ≥0	earlyTerminationAge: 2
minVisiblty4Tracking	Minimum visibility of target bounding box to be considered valid	Float, 0.0 to 1.0	minVisiblty4Tracking: 0.1

2.3 GST-NVSTREAMMUX

The `Gst-nvstreammux` plugin forms a batch of frames from multiple input sources. When connecting a source to `nvstreammux` (the muxer), a new pad must be requested from the muxer using `gst_element_get_request_pad()` and the pad template `"sink_%u"`. For more information, see `link_element_to_streammux_sink_pad()` in the DeepStream app source code.

The muxer forms a batched buffer of batch-size frames. (batch-size is specified using the `gst_object` property.)

If the muxer's output format and input format are the same, the muxer forwards the frames from that source as a part of the muxer's output batched buffer. The frames are returned to the source when muxer gets back its output buffer. If the resolution is not the same, the muxer scales frames from the input into the batched buffer and then returns the input buffers to the upstream component.

The muxer pushes the batch downstream when the batch is filled or the batch formation timeout `batched-pushed-timeout` is reached. The timeout starts running when the first buffer for a new batch is collected.

The muxer uses a round-robin algorithm to collect frames from the sources. It tries to collect an average of $(\text{batch-size}/\text{num-source})$ frames per batch from each source (if all sources are live and their frame rates are all the same). The number varies for each source, though, depending on the sources' frame rates.

The muxer outputs a single resolution (i.e. all frames in the batch have the same resolution). This resolution can be specified using the `width` and `height` properties. The muxer scales all input frames to this resolution. The `enable-padding` property can be set to `true` to preserve the input aspect ratio while scaling by padding with black bands.

For DGPU platforms, the GPU to use for scaling and memory allocations can be specified with the `gpu-id` property.

For each source that needs scaling to the muxer's output resolution, the muxer creates a buffer pool and allocates four buffers each of size:

$$\text{output-width} * \text{output-height} * f$$

Where f is 1.5 for NV12 format, or 4.0 for RGBA. The memory type is determined by the `nvbuf-memory-type` property.

Set the `live-source` property to `true` to inform the muxer that the sources are live. In this case the muxer attaches the PTS of the last copied input buffer to the batched Gst Buffer's PTS. If the property is set to `false`, the muxer calculates timestamps based on the frame rate of the source which first negotiates capabilities with the muxer.

The muxer attaches an `NvDsBatchMeta` metadata structure to the output batched buffer. This meta contains information about the frames copied into the batch (e.g. source ID of the frame, original resolutions of the input frames, original buffer PTS of the input frames). The source connected to the `Sink_N` pad will have `pad_index` `N` in `NvDsBatchMeta`.

The muxer supports addition and deletion of sources at run time. When the muxer receives a buffer from a new source, it sends a `GST_NVEVENT_PAD_ADDED` event. When a muxer sink pad is removed, the muxer sends a `GST_NVEVENT_PAD_DELETED` event. Both events contains the source ID of the source being added or removed (see `sources/includes/gst-nvevent.h`). Downstream elements can reconfigure when they receive these events. Additionally, the muxer also sends a `GST_NVEVENT_STREAM_EOS` to indicate EOS from the source.

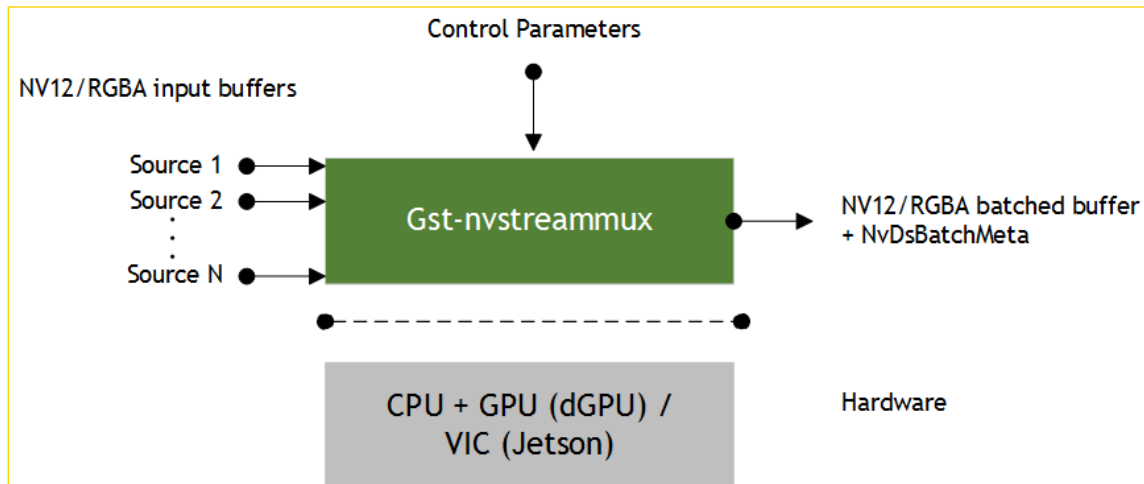


Figure 3. The Gst-nvstreammux plugin

2.3.1 Inputs and Outputs

► Inputs

- NV12/RGBA buffers from an arbitrary number of sources

► Control Parameters

- batch-size
- batched-push-timeout
- width
- height
- enable-padding
- gpu-id (dGPU only)
- live-source
- nvbuf-memory-type

► Output

- NV12/RGBA batched buffer
- GstNvBatchMeta (meta containing information about individual frames in the batched buffer)

2.3.2 Features

Table 9 summarizes the features of the plugin.

Table 9. Features of the Gst-nvstreammux plugin

Feature	Description	Release
Configurable batch size	—	DS 2.0
Configurable batching timeout	—	DS 2.0
Allows multiple input streams with different resolutions	—	DS 2.0
Allows multiple input streams with different frame rates	—	DS 2.0
Scales to user-determined resolution in muxer	—	DS 2.0
Scales while maintaining aspect ratio with padding	—	DS 2.0
Multi-GPU support	—	DS 2.0
Input stream DRC support	—	DS 3.0
User-configurable CUDA memory type (Pinned/Device/Unified) for output buffers	—	DS 3.0
Custom message to inform application of EOS from individual sources	—	DS 3.0
Supports adding and deleting run time sinkpads (input sources) and sending custom events to notify downstream components	—	DS 3.0
Supports RGBA data handling at output	—	DS 3.0

2.3.3 Gst Properties

Table 10 describes the Gst-nvstreammux plugin's Gst properties.

Table 10. Gst-nvstreammux plugin, Gst properties

Property	Meaning	Type and Range	Example Notes
batch-size	Maximum number of frames in a batch.	Integer, 0 to 4,294,967,295	batch-size=30
batched-push-timeout	Timeout in microseconds to wait after the first buffer is available to push the batch even if a complete batch is not formed.	Signed integer, -1 to 2,147,483,647	batched-push-timeout=40000 40 msec
width	If non-zero, muxer scales input frames to this width.	Integer, 0 to 4,294,967,295	width=1280
height	If non-zero, muxer scales input frames to this height.	Integer, 0 to 4,294,967,295	height=720

Property	Meaning	Type and Range	Example Notes
enable-padding	Maintains aspect ratio by padding with black borders when scaling input frames.	Boolean	enable-padding=1
gpu-id	ID of the GPU on which to allocate device or unified memory to be used for copying or scaling buffers. (dGPU only.)	Integer, 0 to 4,294,967,295	gpu-id=1
live-source	Indicates to muxer that sources are live, e.g. live feeds like an RTSP or USB camera.	Boolean	live-source=1
nvbuf-memory-type	Type of memory to be allocated. For dGPU: 0 (nvbuf-mem-default): Default memory, cuda-device 1 (nvbuf-mem-cuda-pinned): Pinned/Host CUDA memory 2 (nvbuf-mem-cuda-device) Device CUDA memory 3 (nvbuf-mem-cuda-unified): Unified CUDA memory For Jetson: 0 (nvbuf-mem-default): Default memory, surface array 4 (nvbuf-mem-surface-array): Surface array memory	Integer, 0-4	nvbuf-memory-type=1

2.4 GST-NVSTREAMDEMUX

The `Gst-nvstreamdemux` plugin demuxes batched frames into individual buffers. It creates a separate Gst Buffer for each frame in the batch. It does not copy the video frames. Each Gst Buffer contains a pointer to the corresponding frame in the batch.

The plugin pushes the unbatched Gst Buffer objects downstream on the pad corresponding to each frame's source. The plugin gets this information through the `NvDsBatchMeta` attached by `Gst-nvstreammux`. The original buffer timestamps (PTS) of individual frames are also attached back to the Gst Buffer.

Since there is no frame copy, the input Gst Buffer is not returned upstream immediately. When all of the non-batched Gst Buffer objects demuxed from an input batched Gst Buffer are returned to the demuxer by the downstream component, the input batched Gst Buffer is returned upstream.

The demuxer does not scale the buffer back to the source's original resolution even if Gst-nvstreammux has scaled the buffers.

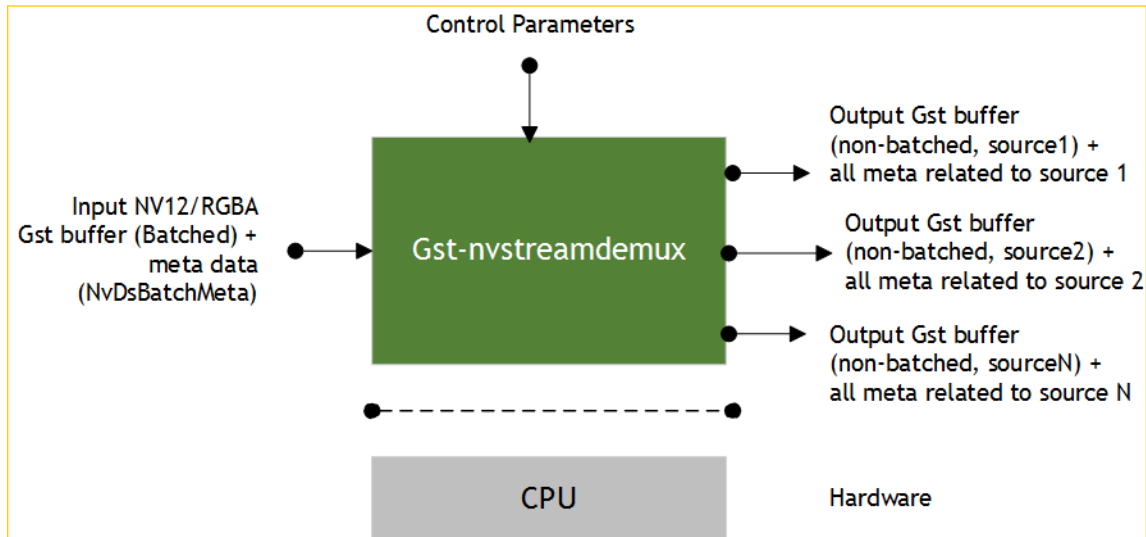


Figure 4. The Gst-nvstreamdemux plugin

2.4.1 Inputs and Outputs

► Inputs

- Gst Buffer (batched)
- NvDsBatchMeta
- Other meta

► Control parameters

- None

► Output

- Gst Buffer (non-batched, single source)
- Meta related to each Gst Buffer source

2.5 GST-NVMULTISTREAMTILER

The `Gst-nvmultistreamtiler` plugin composites a 2D tile from batched buffers. The plugin accepts batched NV12/RGBA data from upstream components. The plugin composites the tile based on stream IDs, obtained from `NvDsBatchMeta` and `NvDsFrameMeta` in row-major order (starting from source 0, left to right across the top row, then across the next row). Each source frame is scaled to the corresponding location in the tiled output buffer. The plugin can reconfigure if a new source is added and it exceeds the space allocated for tiles. It also maintains a cache of old frames to avoid display flicker if one source has a lower frame rate than other sources.

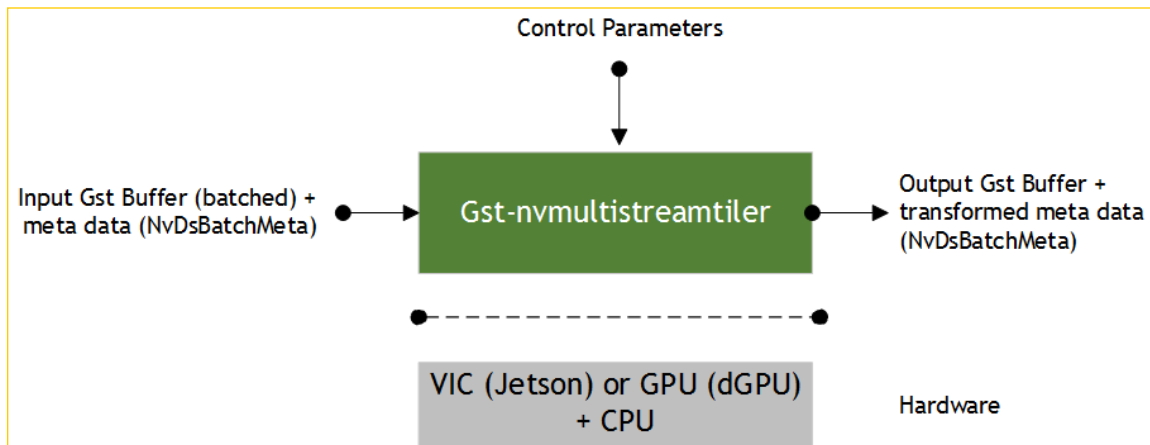


Figure 5. The `Gst-nvmultistreamtiler` plugin

2.5.1 Inputs and Outputs

► Inputs

- `Gst Buffer batched buffer`
- `NvDsBatchMeta` with `Gst Buffer batched` (batch is one or more buffers)

Formats supported: NV12/RGBA

► Control Parameters

- `rows`
- `columns`
- `width`
- `height`
- `gpu-id` (dGPU only)
- `show-source`
- `nvbuf-memory-type`
- `custom-tile-config`

► Output

- Gst Buffer (single frame) with composited input frames
- Transformed metadata (NvDsBatchMeta)

Formats supported: NV12/RGBA

2.5.2 Features

Table 11 summarizes the features of the plugin.

Table 11. Features of the Gst-nvmultistreamtiler plugin

Feature	Description	Release
Composites a 2D tile of input buffers	—	DS 2.0
Scales bounding box with metadata coordinates according to scaling and position in tile	—	DS 2.0
Multi-GPU support	—	DS 2.0
Shows expanded preview for a single source	—	DS 2.0
User configurable CUDA memory type (Pinned/Device/Unified) for output buffers	—	DS 3.0
Reconfigures 2D tile for new sources added at runtime	—	DS 3.0

2.5.3 Gst Properties

Table 12 describes the Gst-nvmultistreamtiler plugin's Gst properties.

Table 12. Gst-nvmultistreamtiler plugin, Gst properties

Property	Meaning	Type and Range	Example Notes
Rows	Number of rows in 2D tiled output	Integer, 1 to 4,294,967,295	row=2
Columns	Number of columns in 2D tiled output	Integer, 1 to 4,294,967,295	columns=2
Width	Width of 2D tiled output in pixels	Integer, 16 to 4,294,967,295	width=1920
Height	Height of 2D tiled output in pixels	Integer, 16 to 4,294,967,295	height=1080

Property	Meaning	Type and Range	Example Notes
show-source	Scale and show frames from a single source. -1: composite and show all sources For values ≥ 0 , frames from that source are zoomed.	Signed integer, -1 to 2,147,483,647	show-source=2
gpu-id	ID of the GPU on which device/unified memory is to be allocated, and in which buffers are copied or scaled. (dGPU only.)	Integer, 0 to 4,294,967,295	gpu-id=1
nvbuf-memory-type	Type of CUDA memory to be allocated. For dGPU: 0 (nvbuf-mem-default): Default memory, cuda-device 1 (nvbuf-mem-cuda-pinned): Pinned/Host CUDA memory 2 (nvbuf-mem-cuda-device) Device CUDA memory 3 (nvbuf-mem-cuda-unified): Unified CUDA memory For Jetson: 0 (nvbuf-mem-default): Default memory, surface array 4 (nvbuf-mem-surface-array): Surface array memory	Integer, 0-4	nvbuf-memory-type=1
custom-tile-config	Custom tile position and resolution. Can be configured programmatically for all or none of the sources.	Values of enum CustomTileConfig	Reserved for future use. Default: null.

2.6 GST-NVDSOSD

This plugin draws bounding boxes, text, and region of interest (**RoI**) polygons. (Polygons are presented as a set of lines.)

The plugin accepts an RGBA buffer with attached metadata from the upstream component. It draws bounding boxes, which may be shaded depending on the